

# 05 – Apache Hadoop

Abdel Dadouche  
DJZ Consulting

[adadouche@hotmail.com](mailto:adadouche@hotmail.com)  
@adadouche

# Hadoop historical facts

# Some historical facts

1997	First appearance of the term “ <b>Big Data</b> ” by NASA to designate the challenges working on large volumes of unstructured data
2001	The 3V's (Volume / Variety / Velocity)
2003 & 2004	Google releases the <b>GFS</b> (Google File System) & <b>MapReduce</b> whitepapers
2004	<a href="#">Doug Cutting</a> & <a href="#">Mike Cafarella</a> release <a href="#">Nutch</a> , a highly extensible and scalable open source web crawler project that implemented the first MapReduce facility on a Distributed File System
2006	Based on Nutch, Doug Cutting (now at Yahoo!) is enhancing the search engine indexation. This was then moved under the new Hadoop subproject.

# Where does the Elephant and Name Hadoop come from?



His son made up the name after his  
toy elephant!



Source: [https://en.wikipedia.org/wiki/Big\\_data](https://en.wikipedia.org/wiki/Big_data)  
<https://twitter.com/theregister/status/463858134683877378>

# Some historical facts

April 2006	Apache Hadoop 0.1.0 was released
2008	<a href="#">Cloudera</a> was founded by three engineers from Google, Yahoo! and Facebook (Christophe Bisciglia, Amr Awadallah and Jeff Hammerbacher)
2008	A Yahoo! Hadoop cluster beat the <a href="#">Terabyte Sort Benchmark</a> (1 TB of data in 209 seconds)
2009	<a href="#">MapR</a> was founded
2010	Dhruba Borthakur <a href="#">claimed</a> that Facebook had the largest Hadoop cluster with 21 PB of storage
2011	<a href="#">HortonWorks</a> was founded
December 2011	Apache Hadoop 1.0.0 was released
2014	Apache Spark 1.0.0 was released
2014	Apache Spark beat the <a href="#">Terabyte Sort Benchmark</a> (100 TB in 1406 seconds)

# **The key principles behind Hadoop**

# The key principles behind the Hadoop platform

- An open & extendable platform
- Distributed architecture
- Provide API and tools
- Easy to install, deploy & maintain on a multi-platform environment
- Scalable & Reliable
- A more agile & flexible approach / philosophy to store and process data !

# The key principles behind Hadoop

An open & extendable platform

- Apache Foundation Open Source project since 2009:

<https://hadoop.apache.org>

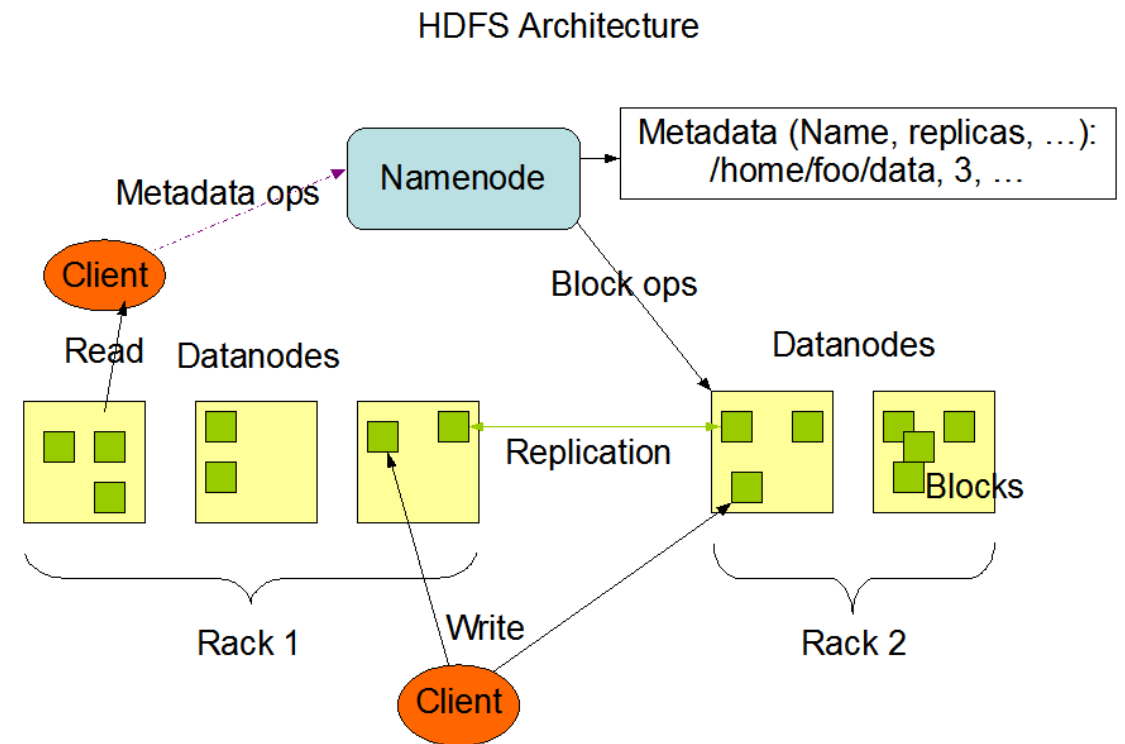
- Developed in Java (multiplatform by design) but can now also use native libraries to improve performances
- Anyone can contribute with code, reporting issues and even start a subproject



# The key principles behind Hadoop

## Distributed Data Storage

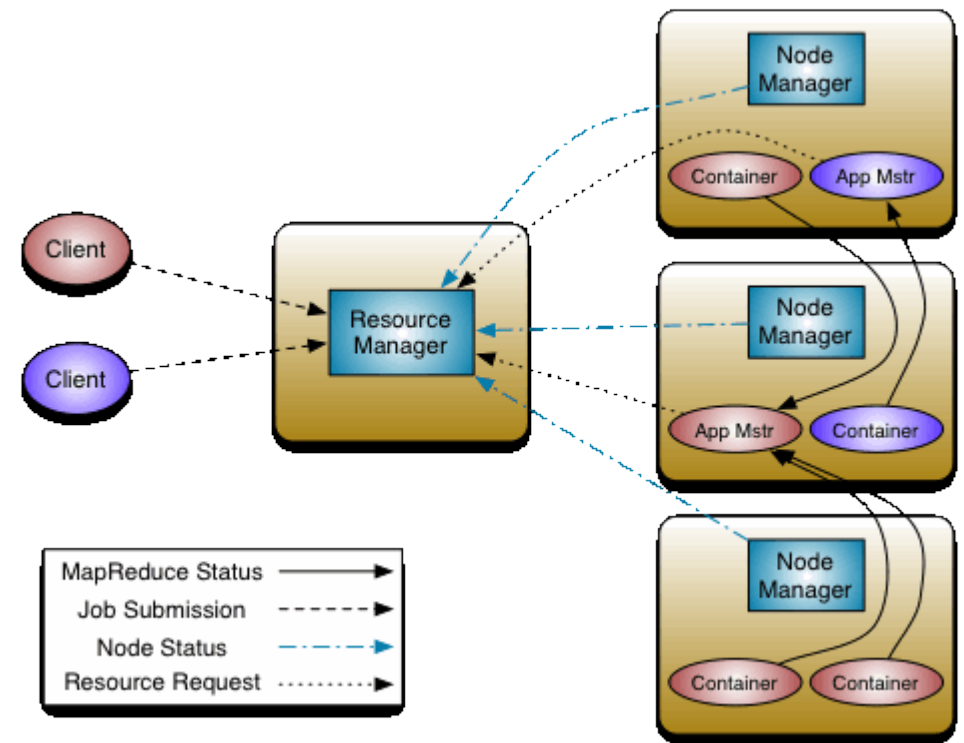
- HDFS use a master/slave architecture
- The NameNode (master) manages the file system namespace and regulates access to files by clients.
- The DataNodes (slave), one per node in the cluster, manage storage attached to the nodes that they run on.



# The key principles behind Hadoop

## Distributed Computing, YARN & MapReduce

- The fundamental of YARN is to split the resource management and job scheduling/monitoring into separate daemons
- The ResourceManager arbitrates resources among all the applications
- The NodeManager is the per-machine agent who is responsible for containers, monitoring their resource usage (cpu, memory, disk, network) and reporting the same to the Resource Manager.



# The key principles behind Hadoop

## API and Tools for Developers

- Multiple API available for HDFS & MapReduce:

Java, C++, Python, NodeJS...

- Multiple tools available to administer and monitor your Hadoop clusters:

Command line tools (fs, job, dfsadmin, jobtracker...)

Ambari, ZooKeeper and many vendor specific tools...

### Warning:

There are plenty of duplicate APIs.

Some are not maintained or simply outdated

# The key principles behind Hadoop

Easy to install, deploy & maintain on a multi-platform environment

- It's written in Java!
- Can be downloaded and installed from :  
  
<https://hadoop.apache.org/releases.html>
- Also available for most Linux distribution as packages (via apt-get, yum, rpm etc.)
- Can be installed on Windows as well! (this will be your first assignment)

# The key principles behind Hadoop

Reliable, Scalable & Highly Available

- Data is chunked then distributed across the cluster
  - Multiple copies of the same chunk exists on different nodes
- New nodes can be added without downtime
  - Data chunks will start being added
- If a node crashes, it won't affect the availability of the all data
- Job logic (code) is distributed across the cluster
- Each node will process the data chunks locally stored
- Each data chunks will be processed by only one node (shared-nothing architecture)
- If a node crashes, the job can be executed on a different node where data chunk is available

# The key principles behind Hadoop

## Schema on Read vs Schema on Write

- Schema on Write (RDBMS)
  - Need to build a schema before bringing the data
  - Data must fit the schema before being accessed
  - Schema must be updated before any new data is loaded
- Benefits
  - Fast read
  - Governance
- Schema on Read (Hadoop)
  - Data loaded directly to the file system
  - No transformation is needed
  - Late binding when specific column are read using SerDE (Serializer/Deserializer)
- Benefits
  - Fast load
  - Agility & Flexibility

# Conclusion

- The Hadoop platform will allow you to build your Large Scale Processing applications thanks to :
  - Distributed data storage
  - Distributed data processing
  - Fault Tolerance
  - Linear scalability
  - Hardware agnostic
  - And so much more..
- The platform / ecosystem is huge!

# The Hadoop Core



# The core subproject of Apache Hadoop

- Hadoop Common :
  - Utility services used by the other components
- Hadoop Distributed File System (HDFS)
- Hadoop MapReduce v2
- Hadoop YARN : Yet Another Resource Negotiator
  - Resource management and tasks scheduling

Apache Hadoop <https://hadoop.apache.org/>

hadoop

23,167 commits

276 branches

325 releases

236 contributors

Apache-2.0

Branch: trunk

New pull request

Find file

Clone or download

hadoop-common-project

hadoop-hdfs-project

hadoop-mapreduce-project

hadoop-yarn-project

# **An Overview of the Hadoop Ecosystem**

# Hadoop Ecosystem

## Data Visualization

SAS Visual Analytics

Tableau

Qlik

SAP Lumira

R

D3.JS

iCharts

Timeline JS

Apache Zeppelin

## System Deployment

Apache Ambari

Apache Mesos

Marathon

Hortonworks HOYA

Apache Bigtop

Deploop

Apache Eagle

Cloudera HUE

Myriad

Brooklyn

Apache Helix

Buildoop

SequenceIQ Cloudbreak

## Data Ingestion

Apache Flume

Apache Sqoop

Facebook Scribe

Apache Chukwa

Apache Kafka

Netflix Suro

Apache Samza

Cloudera Morphline

HIHO

Apache NiFi

Apache ManifoldCF

## Service Programming

Apache Thrift

Apache Zookeeper

Apache Avro

Apache Curator

Apache Karaf

Twitter Elephant Bird

LinkedIn Norbert

## Scheduling & DR

Apache Oozie

LinkedIn Azkaban

Apache Falcon

Shedoscope

## Security

Apache Sentry

Apache Knox Gateway

Apache Ranger

## Frameworks

Jumbune

Spring XD

Cask Data App Platform

## Metadata

Metascope

Apache Tika

## Machine Learning

Apache Mahout

WEKA

Cloudera Oryx

Deeplearning4j

MADlib

H2O

Sparkling Water

Apache SystemML

## Distributed Programming

Apache Ignite

Apache MapReduce

Apache Pig

JAQL

Apache Spark

Apache Storm

Apache Flink

Apache Apex

Netflix PigPen

AMPLAB SIMR

Facebook Corona

Apache REEF

Apache Twill

Damballa Parkour

Apache Hama

Datasalt Pangool

Apache Tez

Apache DataFu

Kangaroo

TinkerPop

Pachyderm MapReduce

Apache Beam

## SQL on Hadoop

Apache Hive

Apache HCatalog

Apache Trafodion

Apache HAWQ

Apache Drill

Cloudera Impala

Facebook Presto

Datasalt Splout SQL

Apache Tajo

Apache Phoenix

Apache MRQL

Kylin

## NoSQL Databases

### Wide Column

Apache HBase

Apache Cassandra

Hypertable

Apache Accumulo

Apache Kudu

Apache Parquet

### Document

MongoDB

RethinkDB

ArangoDB

CouchDB

DynamoDB

Gemfire

### Key-Value

Redis

LinkedIn Voldemort

RocksDB

OpenTSDB

### Graph

Giraph

Neo4j

TitanDB

OrientDB

### Stream Data Model

EventStore

## NewSQL Databases

TokuDB

HandlerSocket

Akiban Server

Drizzle

Haeinsa

SenseiDB

Sky

BayesDB

InfluxDB

VoltDB

SAP HANA

## Distributed File System

Apache HDFS

Red Hat GlusterFS

Quantcast File System

Ceph File System

Lustre File System

Alluxio

GridGain

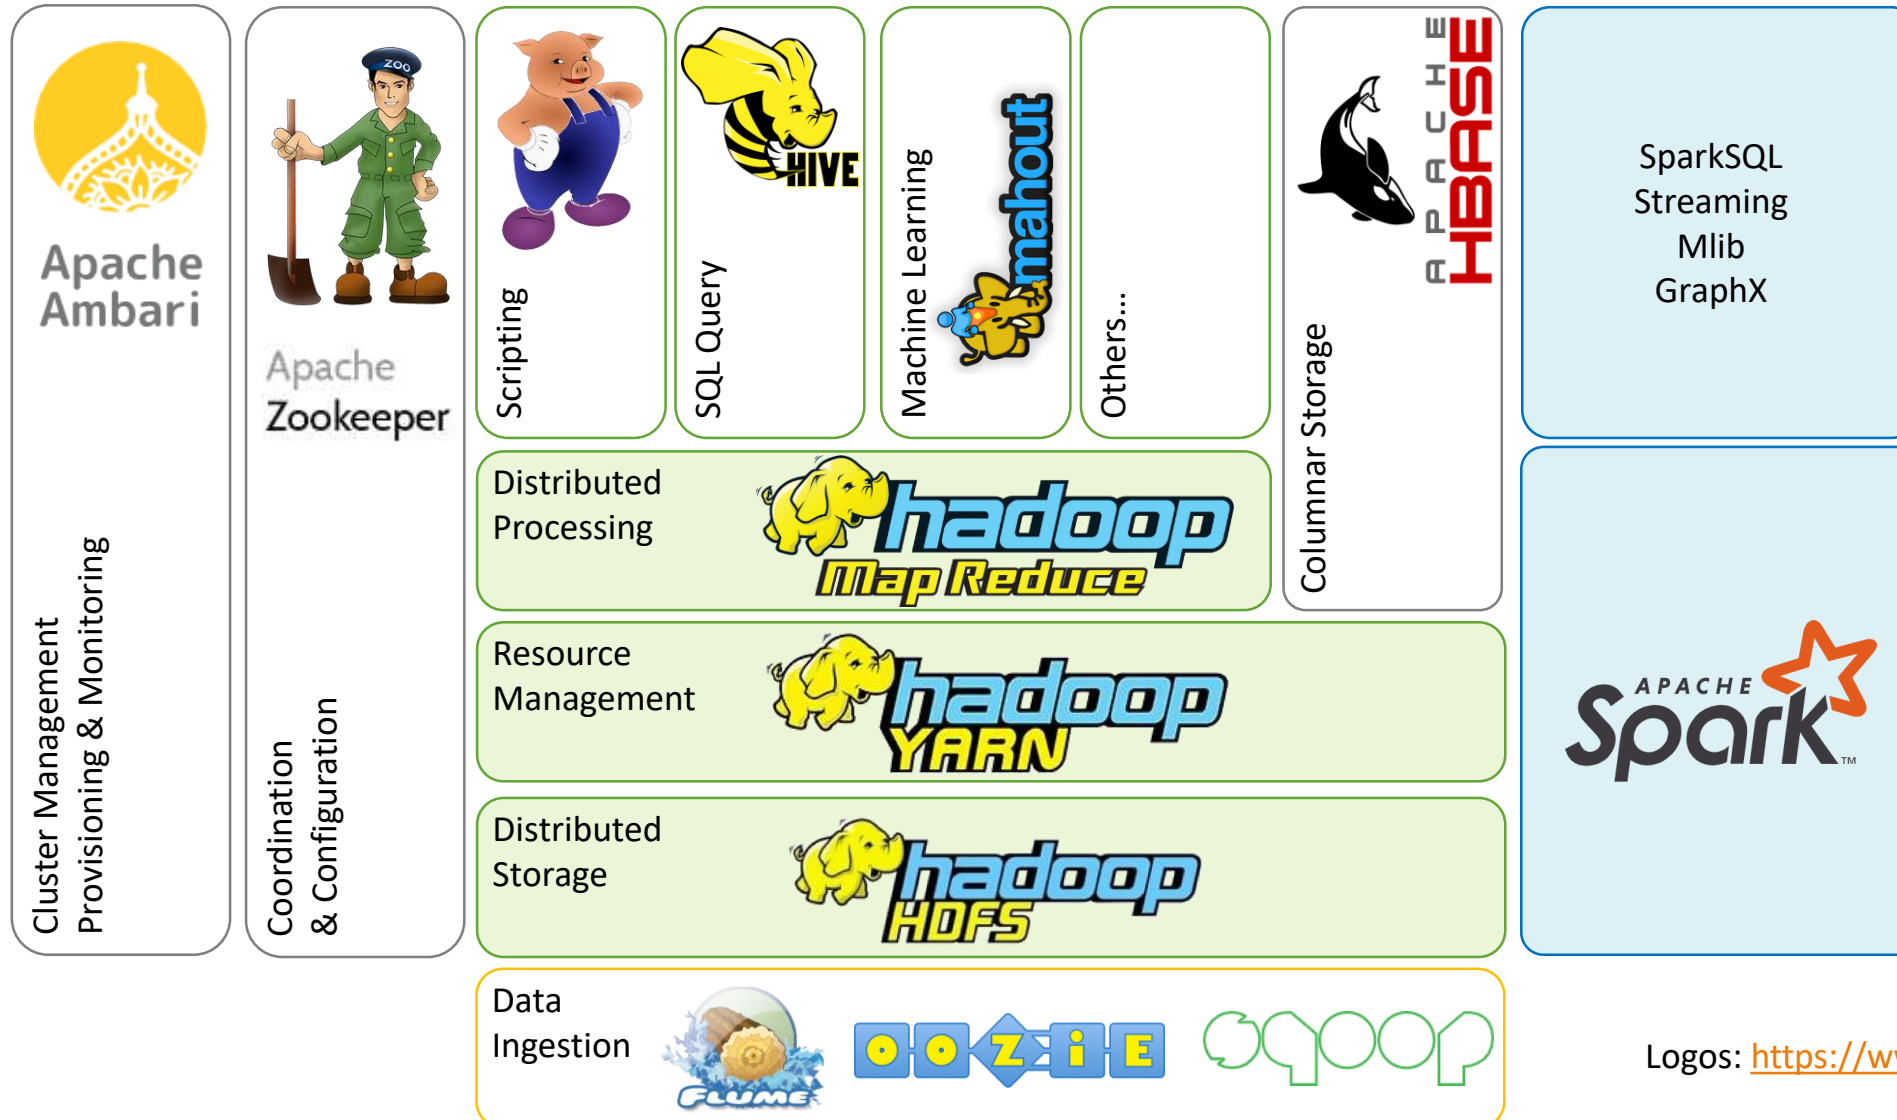
XtreemFS

# The role of ecosystem

- Data visualization
  - System Deployment & Management
  - Data Ingestion/Extraction
  - Service Programming
  - Scheduling
  - Security
  - Metadata extraction
  - Machine Learning
  - Distributed Programming
- SQL Querying
  - Distributed Data Storage
  - Database Models
    - Wide Column
    - Key-Value
    - Document
    - Graph
    - Stream
    - ...

# A Traditional Installation

# A Simple View (of what you will try to use)



# A Traditional Installation

- Apache Ambari
  - Enables system administrators to provision, manage and monitor a Hadoop cluster
- Apache ZooKeeper
  - provide a distributed configuration & synchronization service and naming registry
- Apache Pig
  - high-level language for expressing data analysis programs
- Apache Hive
  - a SQL-like interface to query and analyze data
- Apache Mahout
  - scalable machine learning algorithms
- Apache HBase
  - Bigtable-like capabilities on top of Hadoop
- Apache sqoop
  - Bulk data transfer with structured datastores such as relational databases
- Apache oozie
  - Workflow scheduler system to manage Apache Hadoop jobs
- Apache Flume
  - Collecting, aggregating, and moving large amounts of log data based on streaming data flows

# A Traditional Installation

- Apache Spark
  - Uses resilient distributed dataset (RDD), a distributed read-only multiset of data
  - Spark Core:
    - provides distributed task dispatching, scheduling, and basic I/O functionalities, exposed through an API
  - Spark SQL
    - provides a domain-specific language (DSL) to manipulate DataFrames & SQL language support
  - Spark Streaming
    - ingests data in mini-batches and performs RDD transformations
  - MLlib Machine Learning Library
  - GraphX
    - distributed graph-processing framework



# **The Traditional Data Processing Approaches**

# Data Processing approaches

- **Batch Processing**

- **Input:** Analyze all available data at a certain point in time
- **Results:** The results is available at the end of the process
- **Latency:** The processing time can be long (minutes to hours or days)

- **Micro-batch**

- **Input:** Analyze all available data at a certain point in time every x seconds (smaller files or chunks)
- **Results:** The results is available at the end of the process
- **Latency:** The processing time is short (in seconds)

- **Real Time**

- **Input:** Analyze every new data made available (streams)
- **Results:** The results is available for each new data
- **Latency:** The processing time is really short (in milliseconds)

# The Hadoop Distributions

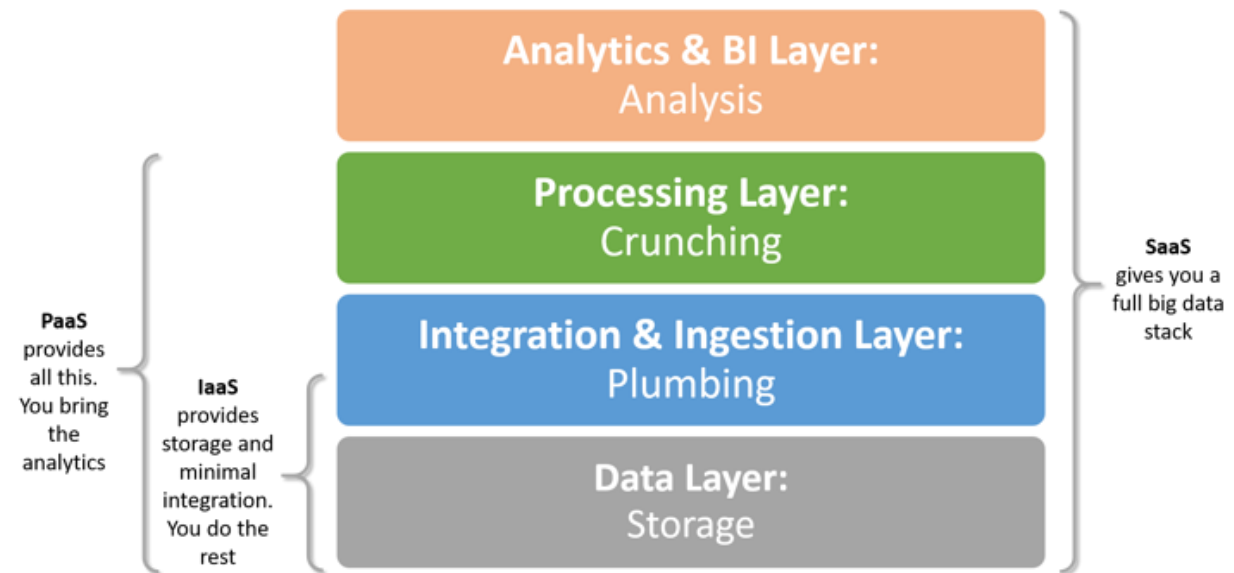
# Open Source, but not only!

- Hortonworks
  - Complete open source Apache Hadoop distribution without additional proprietary software
- Cloudera
  - The core Apache Hadoop distribution with proprietary tools to automate the installation process and other services (monitoring)
- MapR
  - HDFS was replaced by MapRFS to enable more efficient management of data, reliability and ease of use

But they all contribute back to the open source project!

# Also In the Cloud

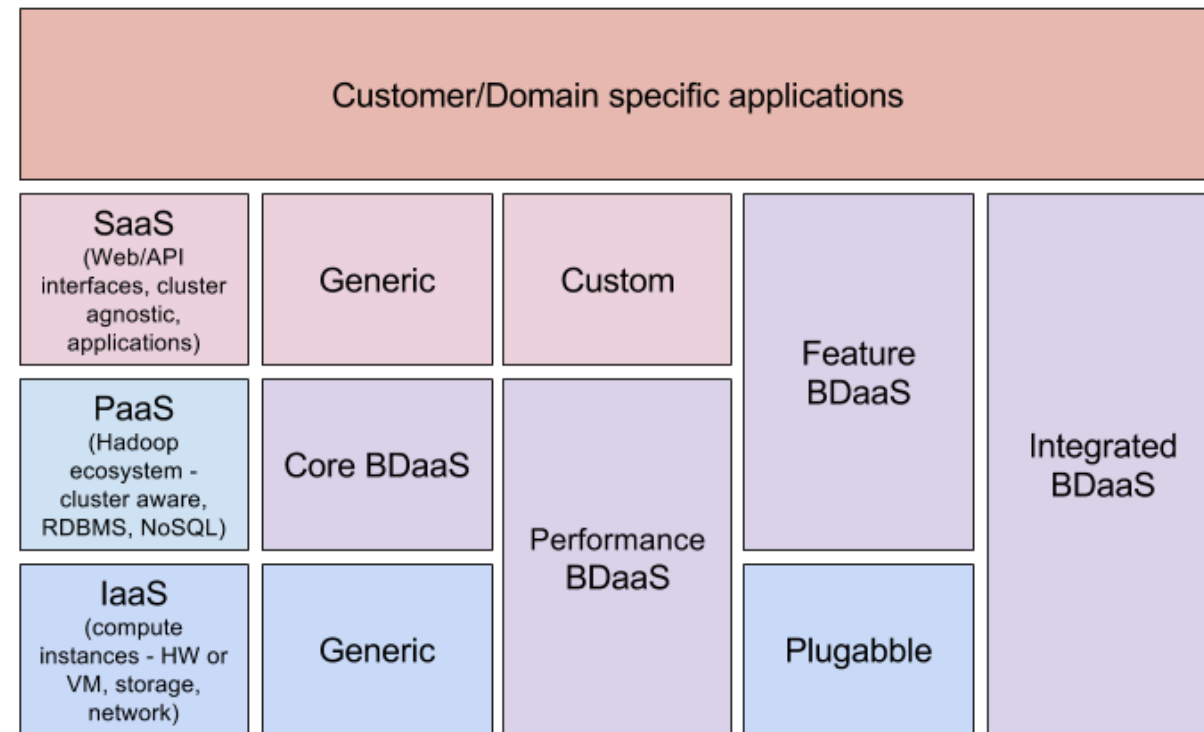
- IaaS : Just compute!
  - Like on-premise but with someone else compute
  - Example: AWS EC2 & EBS
- PaaS : Cloud services to manage your cluster
  - The deployment is fully managed, you don't manage compute anymore
  - Example: Azure HDInsight, AWS EMR & Redshift
- SaaS: The full stack (Big Data + Analytics)
  - You don't care how your data is stored or processed
  - You just use API or tools to access it
  - Examples: Snowflake, Looker



Source: <https://hub.packtpub.com/big-data-as-a-service-bdaas-solutions-comparing-iaas-paas-and-saas/>

# What About BDaaS? A new emerging trend

- Core BDaaS :
  - implement the minimal Hadoop core & a few popular services
- Performance BDaaS:
  - Includes an optimized infrastructure with specifically build hardware
  - Example: Altiscale
- Feature BDaaS
  - Include features beyond the Hadoop core.
  - Example: Qubole
- Integrated BDaaS
  - Combines both the performance and feature





# ODPi – Open Data Platform Initiative

# The Context

- The Apache Hadoop framework and ecosystem has grown really fast:
  - A lot of innovation (streaming, machine learning...)
  - But also a lot of overlap
  - Some frameworks had a short life
- This is causing a slow down in the adoption!
  - People don't know what they should use for their use case
  - People don't know if the framework will be maintained or evolve



# ODPi: Toward Simplification & Standardization?

- The Open Data Platform initiative is a nonprofit organization (but driven by industry/vendors veteran)
- Committed at creating and standardizing big data solutions that work across platforms, systems and products to:
  - Reduce redundancy, overlap and complexity
  - Ease implementation
- Uses a common reference specification called: ODPi Core

Source: <https://www.odpi.org/>

# Hadoop : ODPi Runtime Compliant

- In June 2016, Apache Hadoop, version 2.7, has become ODPi compliant
- ODPi complements and reinforce the Apache Software Foundation work
- The validation test suite guarantees that a distribution of Apache Hadoop complies with the ODPi-defined specifications.
- The test framework and self-certification aligns closely with the Apache Software Foundation by leveraging Apache BigTop for comprehensive packaging, testing, and configuration.

Source: <https://www.odpi.org/announcements/2016/06/27/apache-hadoop-distributions-now-odpi-runtime-compliant>  
<https://github.com/odpi/specs/blob/master/ODPi-Runtime.md>

# Summary

- Hadoop is relatively young but yet mature
- Key principles: distributed, open, extendable, scalable & reliable
- The Hadoop Core: HDFS, MapReduce and YARN
- The type of data processing
- The Hadoop distributions, on-premise, in the cloud, IaaS, PaaS, SaaS & BDaaS
- The ODPi