

07 – Hadoop MapReduce

Abdel Dadouche
DJZ Consulting

adadouche@hotmail.com
@adadouche

Context

- In order to implement a large scale distributed processing system, you need be able to:
 - Split your process into smaller pieces
 - Apply these smaller processing unit on fragment of your data
 - Reconcile the overall process and result
 - Generalize this approach to all type of problems

MapReduce = Map + Reduce operation

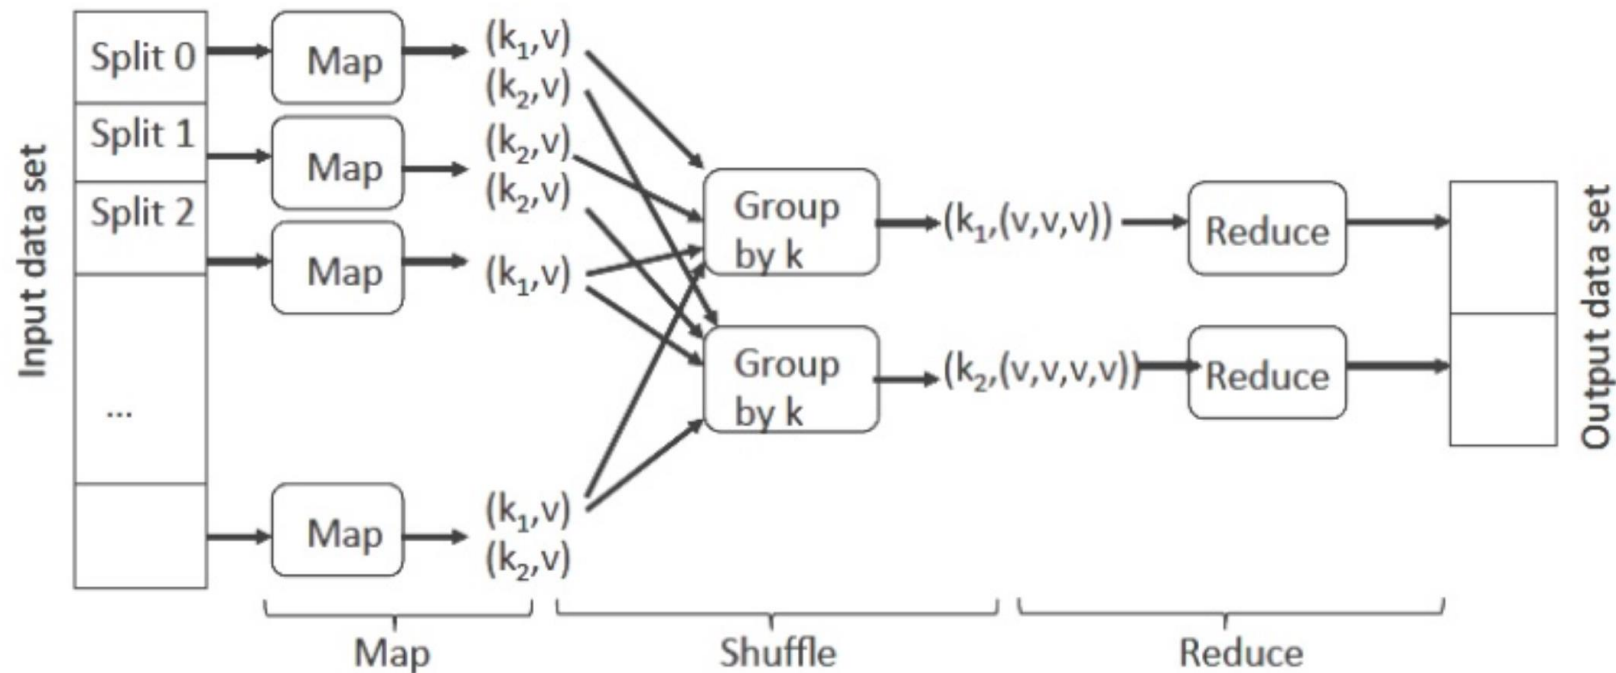
- Map

- Will transform your data into a key/value pairs where the key will be used to group entries and must be chosen based on the problem to solve
- The Map operation must be able to run on fragments of the data (data blocks)
- The execution will take place on the node where the data block is located

- Reduce

- Process each set of values for each keys produced by the Map
- Each set of values is processed by a single node
- There will be one unique result for each set of values (each key) at the end

Here is how to model your problem in 3 ops!



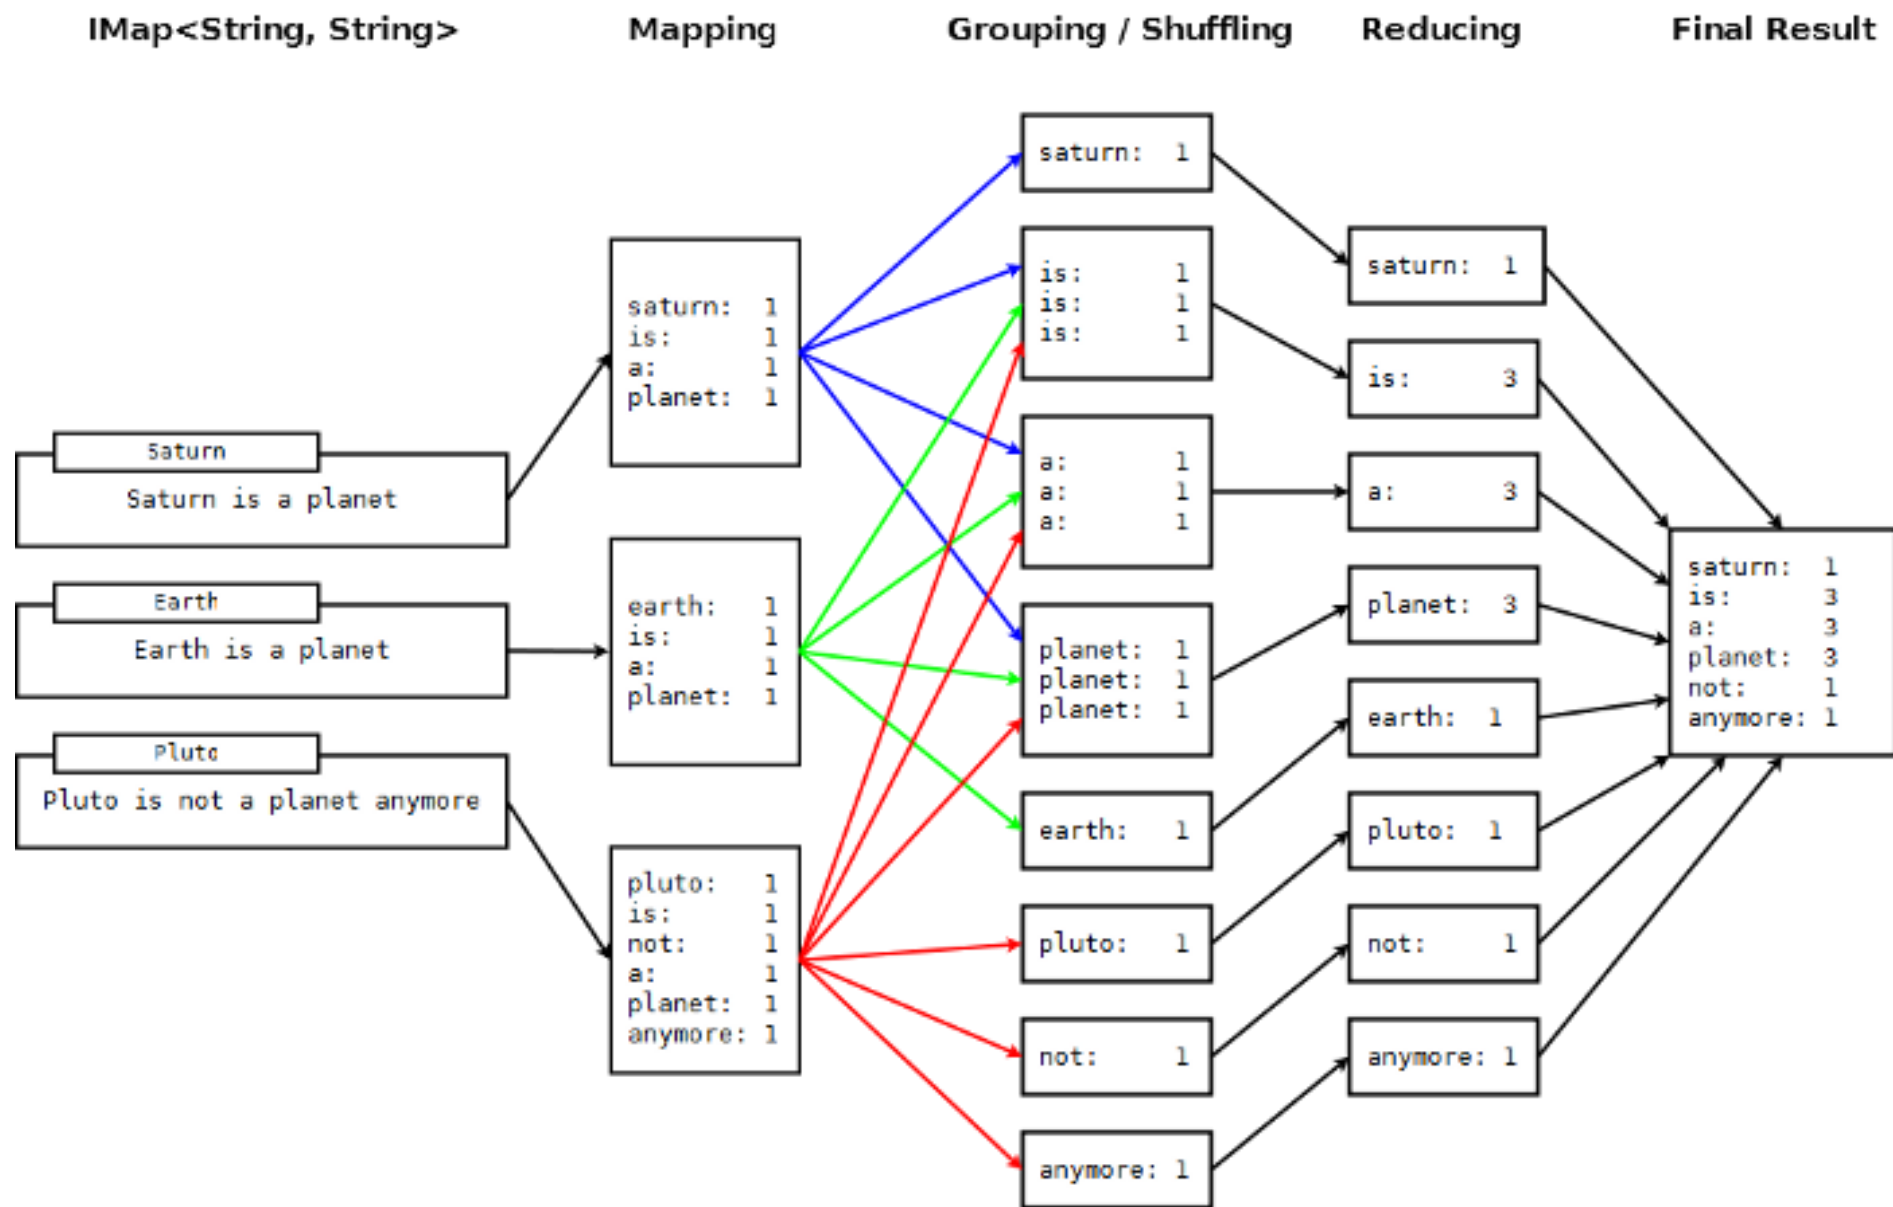
- Split:
 - Represent a data fragment (a data block stored by a data node)
- Map
 - Transform the split into key/value pairs
- Shuffle (optional):
 - Group back the key from each Split
- Reduce
 - Process each set of key/value pairs into a final value

The Word Count example

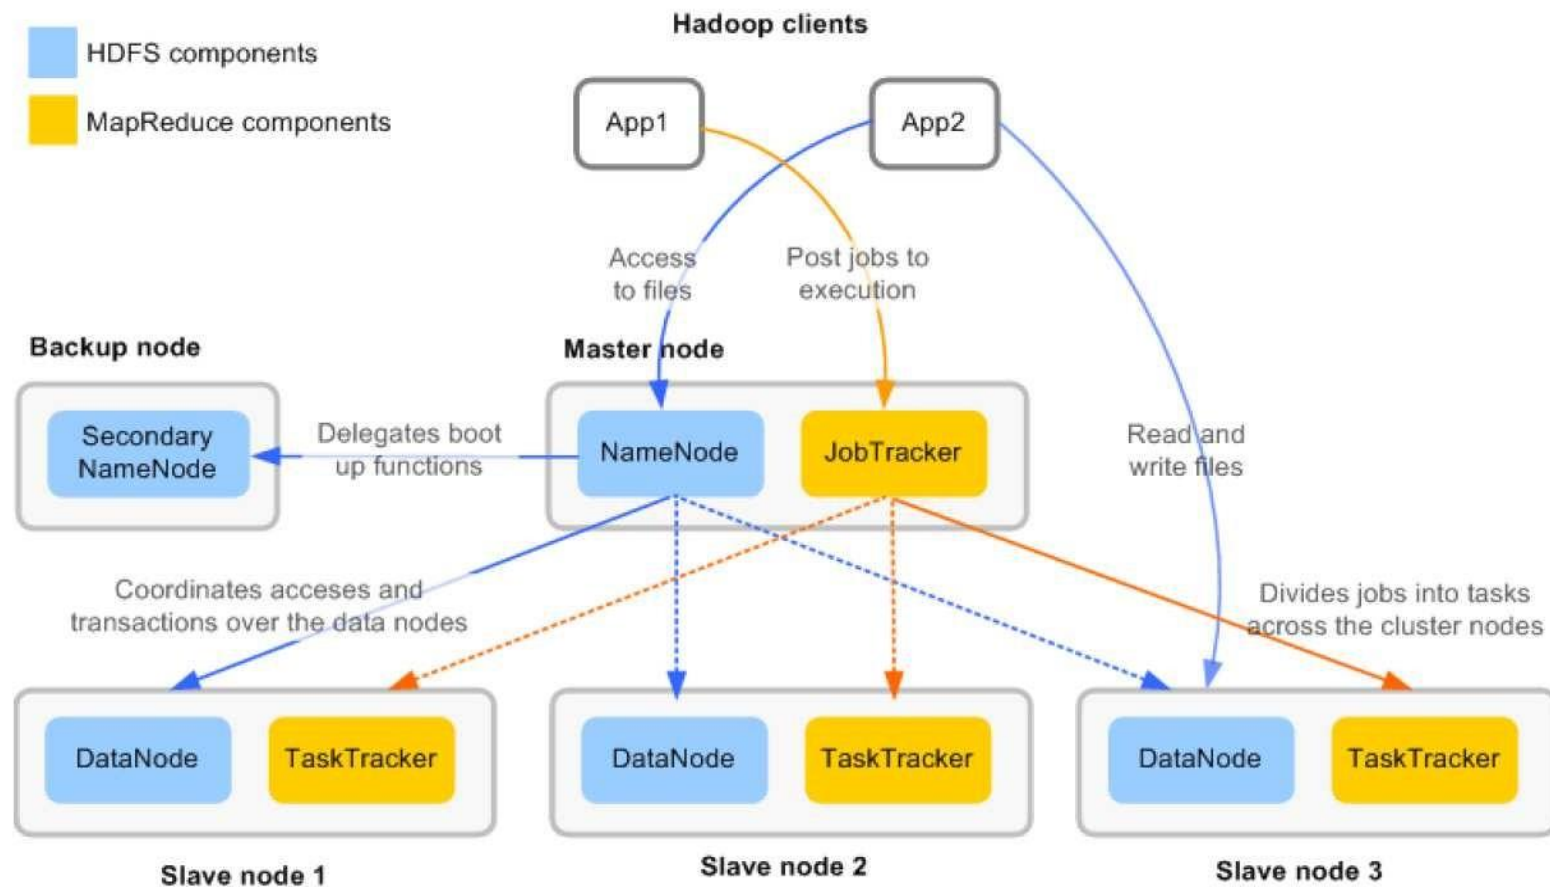
- A record will correspond to a line from the text file separated by a “\n”
- Split:
 - You can choose to have a split per row or for multiple rows
 - As long as each record is stored in a unique split
- Map
 - The key will be the word itself. Every time the word is found, the value (“count”) will be incremented
- Shuffle
 - Group by “word” the “count” from each Map
- Reduce
 - Sum the “count” for each group of identical “word” from the Shuffle

The Word Count example

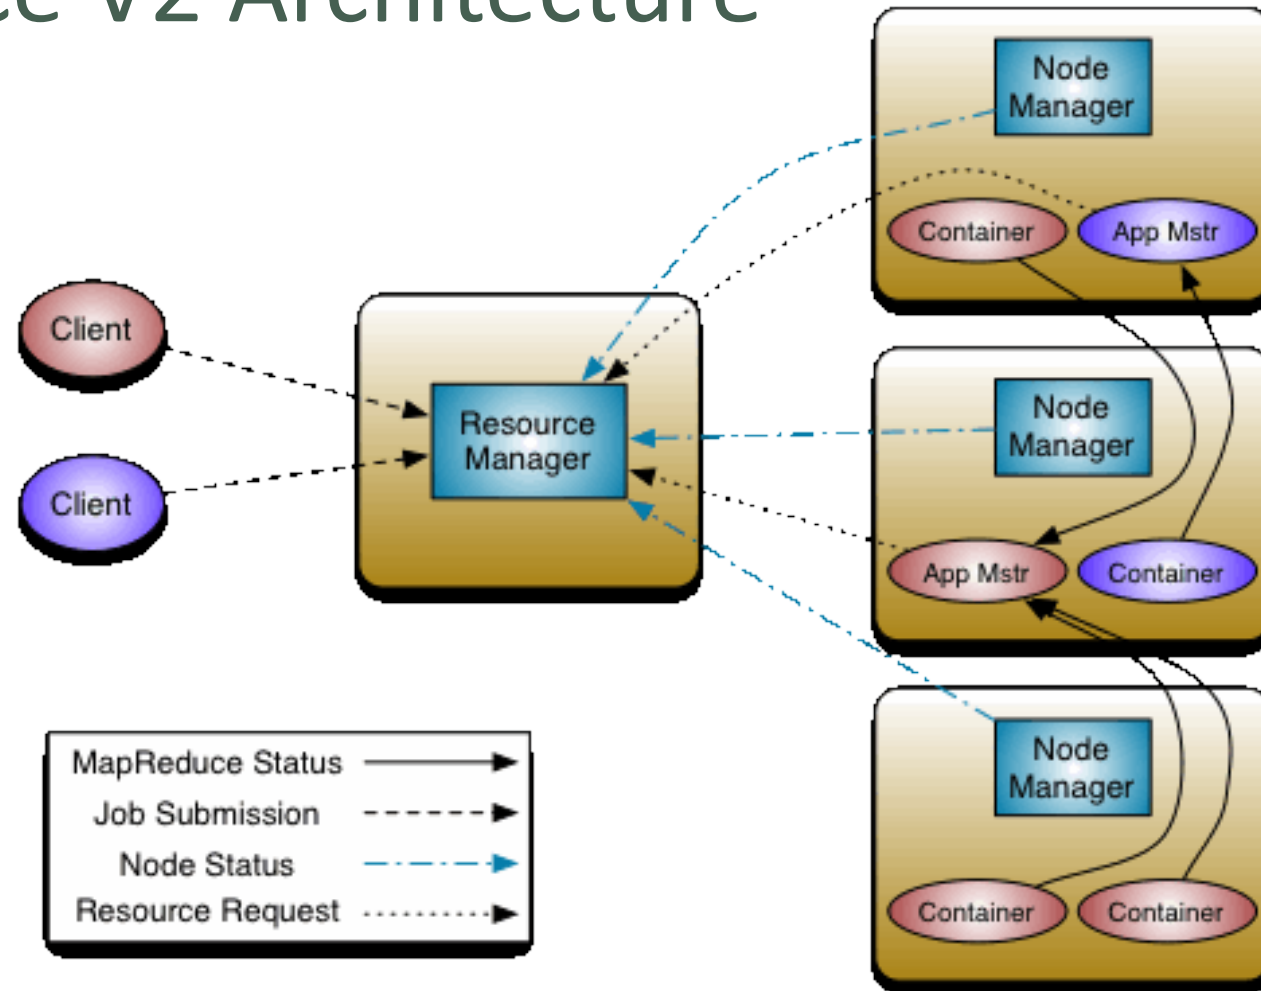
- On a small file or a single node cluster, the performance gain will be fairly limited
- Now, imagine a cluster which will run 5 Mapper and 4 Reducer with 10 Split
 - The Map operation should take only a fifth of the time
 - The Reduce operation should take only a fourth of the time
 - However, you will need to consider the transfer in between nodes (if any)



MapReduce V1 Architecture



MapReduce V2 Architecture



Recap

- Only a few ops: Split, Map, Shuffle, Reduce
- Divide and conquer approach
- Be ready to code in Java!
- The architecture has evolved a lot (with YARN) for more flexibility and better resources allocation & distribution