

12 - Spark SQL - Catalyst Optimizer

Abdel Dadouche
DJZ Consulting

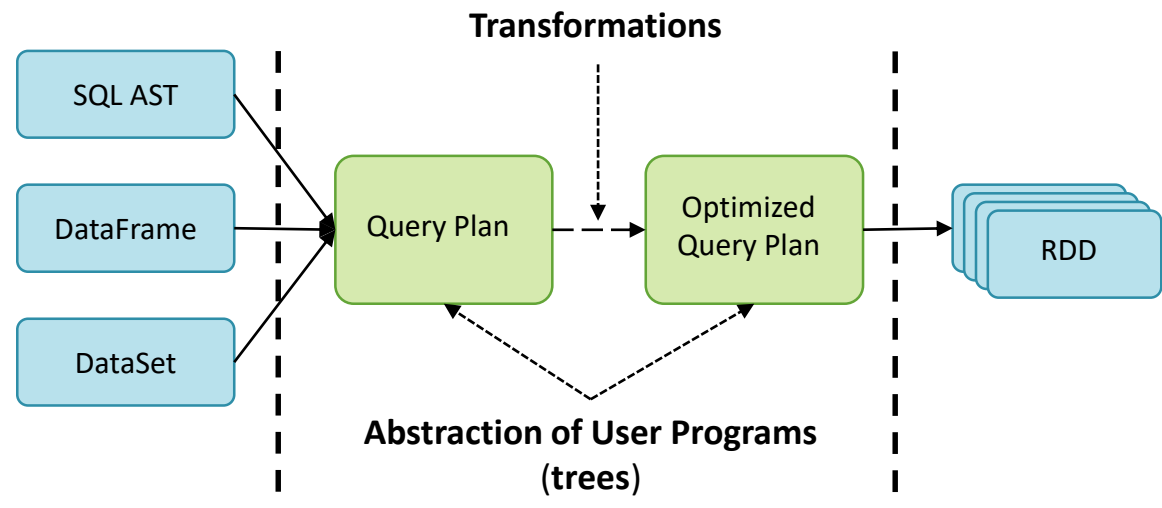
adadouche@hotmail.com
@adadouche

The Catalyst Optimizer

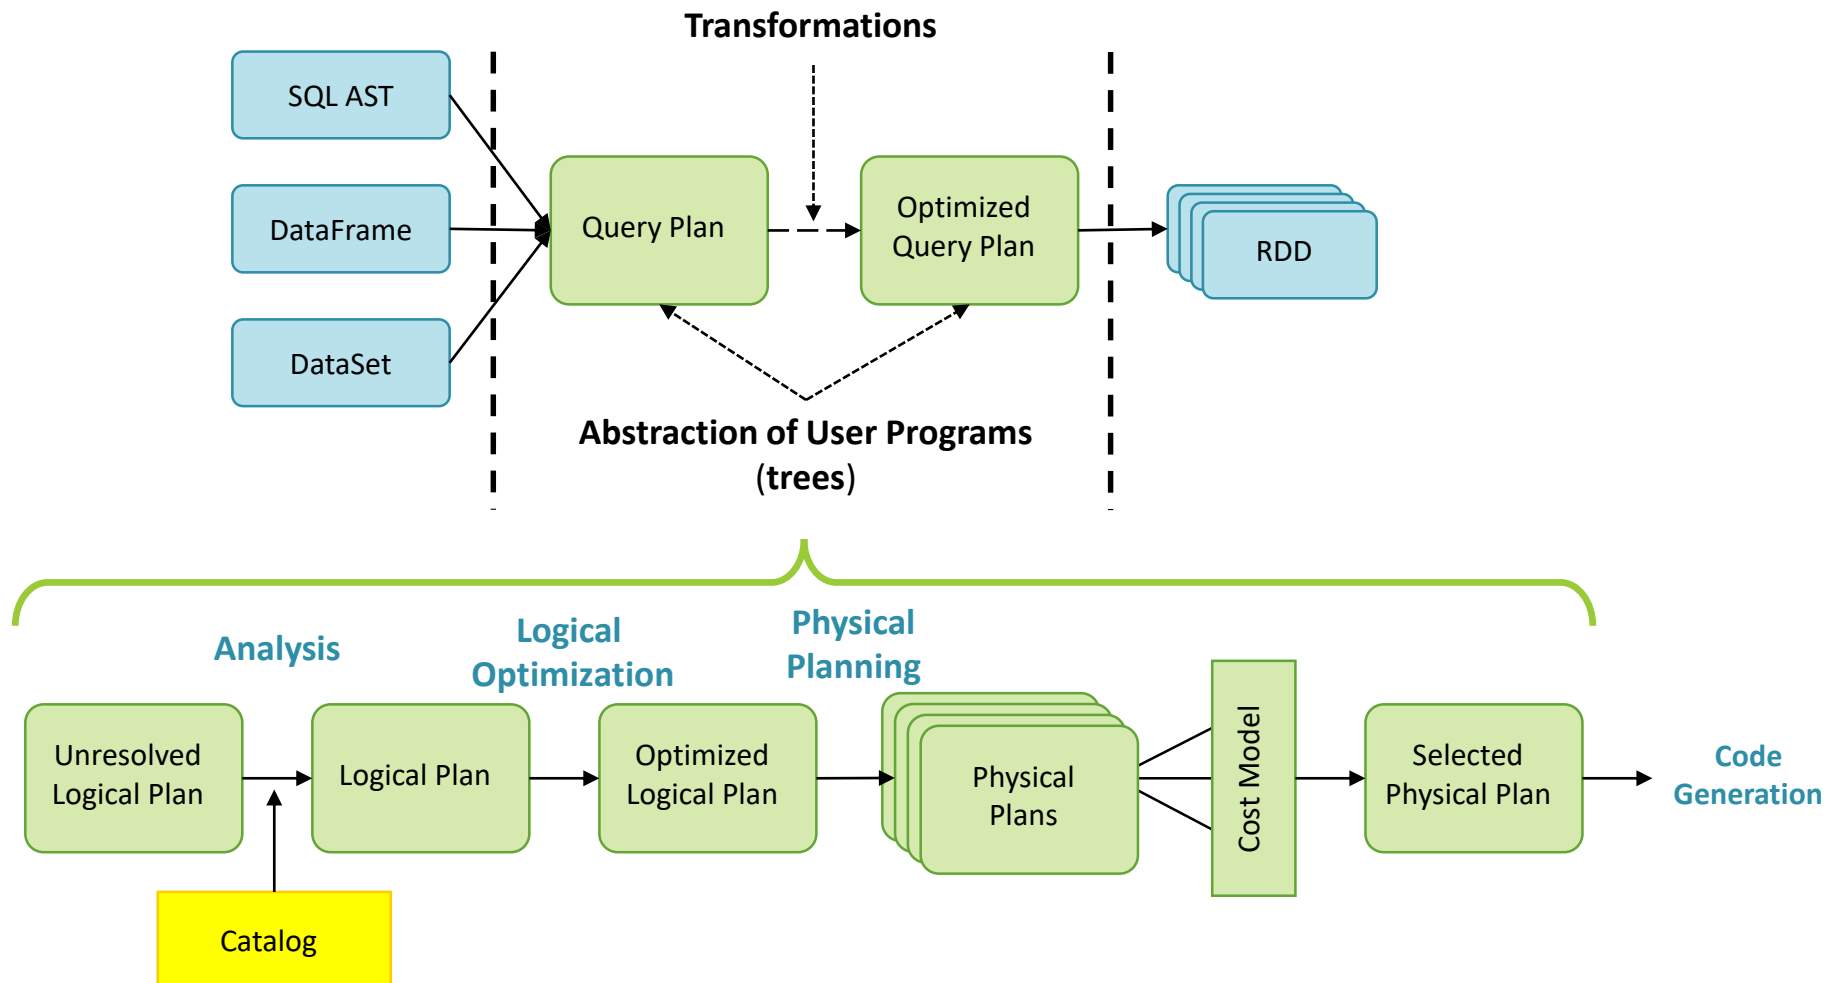
- Initially based on the functional programming constructs in Scala
- Designed around these two key purposes:
 - Easily add new optimization techniques and features to Spark SQL
 - Enable external developers to extend the optimizer
- Offers several public extension points, including external data sources and user-defined types.
- Supports both rule-based and cost-based optimization

Overview of the Catalyst Optimizer

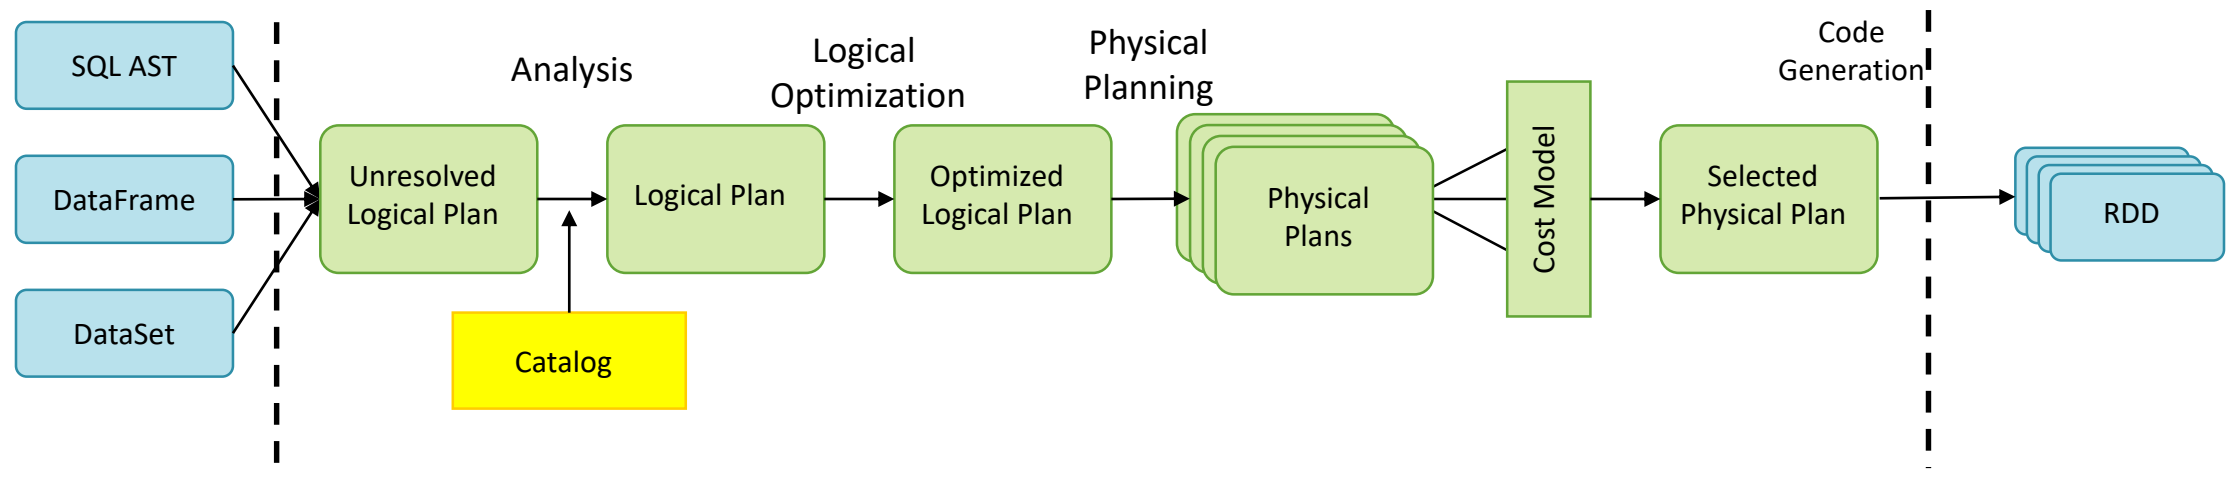
- Can process SQL AST, DataFrame or DataSet
- Apply a series of **transformations** on an initial “Query Plan” to find the “Optimized Query Plan”
- Uses a transformation framework based on a **tree** representation and **rules** for the manipulations
- The output is “code” that will build the RDD stack representing the optimized query plan



Catalyst Optimizer in more details



Catalyst Optimizer in more details



- The Catalyst Optimizer process can be decomposed in 4 **transformations** :
 - analyzing an unresolved logical plan to resolve references
 - logical plan optimization
 - physical planning
 - code generation

Transformations

A Query Example – Attributes & Expressions

```
SELECT    sum(v) as sum
FROM (
  SELECT  t1.id as id,
          1 + 2 + t1.value as v
  FROM    t1 JOIN t2
  WHERE   t1.id = t2.id
  AND     t2.id > 50000
) tmp
```

Attributes

Represents a column of a dataset (t1.id) or a column generated by a specific data operation (the v column)

Expressions :

Represents a new value computed out of an input value (1 + 2 + t1.value)

Attributes & Expressions are represented as trees/nodes

A Query Example – Logical Plan

SELECT sum(v) as sum

FROM (

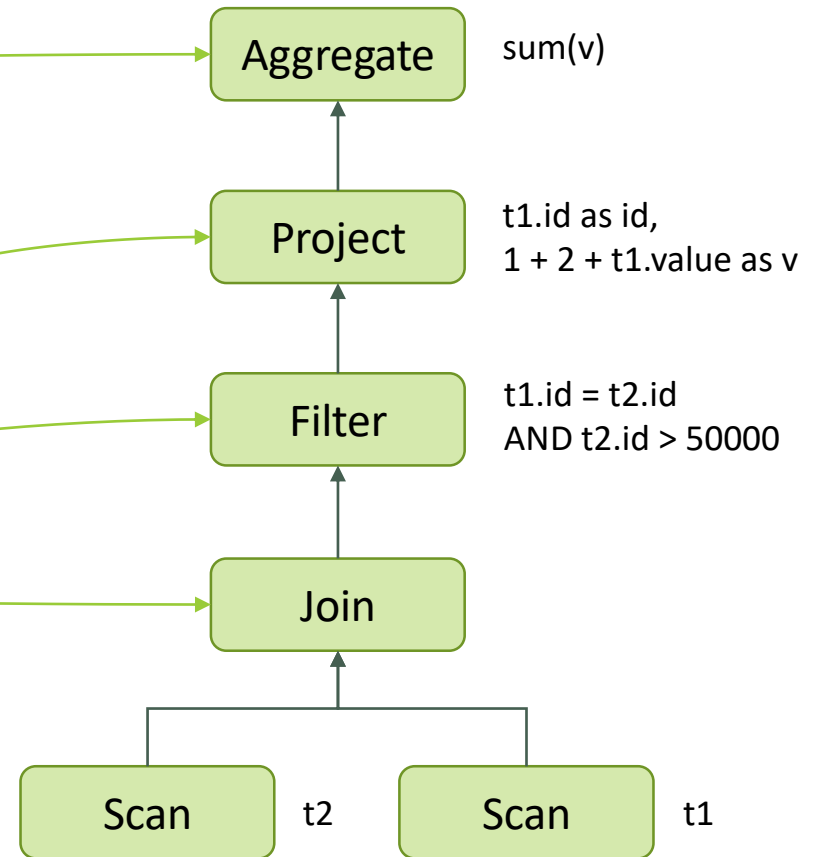
SELECT t1.id as id,
1 + 2 + t1.value as v

FROM t1 JOIN t2

WHERE t1.id = t2.id

AND t2.id > 50000

) tmp



Transformations

- Use abstraction of user programs represented as trees to apply rules
- 2 types of trees transformations
 - Doesn't change the tree type
 - Expression → Expression
 - Logical Plan → Logical Plan
 - Physical Plan → Physical Plan
 - Change the tree type
 - Logical Plan → Physical Plan


Transformations – Trees

- Trees are composed of “node” objects
- Each node has a node type/class
- Each node has zero or more “child” nodes
- New node types are defined in Scala as subclasses of the `TreeNode` class
- These objects are immutable and can only be manipulated using functional transformations

Example : $x + (1 + 2)$

- 3 node classes:
 - `Literal(value: Int)`: a constant value
 - `Attribute(name: String)`: an attribute from an input row, e.g., "x"
 - `Add(left: TreeNode, right: TreeNode)`: sum of two expressions
- In Scala code:

```
Add(Attribute(x), Add(Literal(1), Literal(2)))
```



Transformations – Rules / Transform

- Rules are used to manipulate Trees (a functions from a Tree to another Tree)
- While a rule can run any arbitrary code on its input tree, the most common approach is to use a set of pattern matching functions that find and replace subtrees with a specific structure
- Pattern matching allows extracting values from potentially nested structures.
- In Catalyst, trees offer a **transform** method that applies a pattern matching function recursively on all nodes of the tree, transforming the ones that match each pattern to a result.

Example : $x+(1+2)$

- implement a rule that folds Add operations between constants :

```
tree.transform {  
  case Add(Literal(c1), Literal(c2)) => Literal(c1+c2)  
}
```

- Rules can match multiple patterns in the same transform call, making it very concise to implement multiple transformations at once:


```
tree.transform {  
  case Add(Literal(c1), Literal(c2)) => Literal(c1+c2)  
  case Add(left, Literal(0)) => left  
  case Add(Literal(0), right) => right  
}
```

Transform – Single Rule Example


$x + (1 + 2)$



$x + 3$




Will evaluate $1 + 2$ for every row → waste of calculation resource



Will evaluate $1 + 2$ once → more efficient

Transform – Multiple Rule Example – Batch



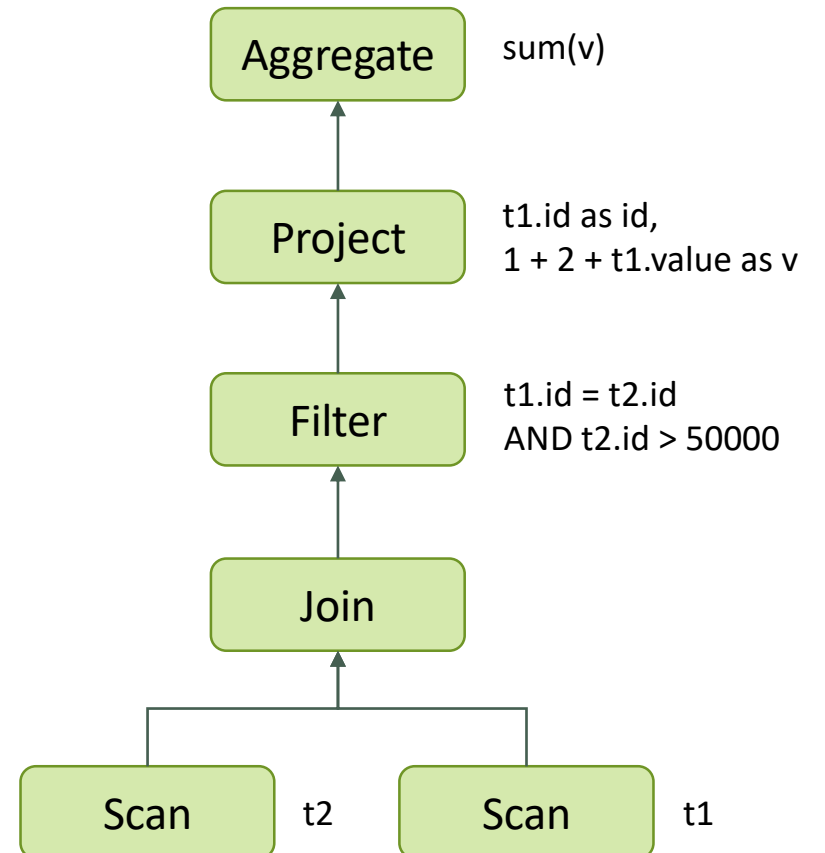
Catalyst Optimizer Phase Analysis

Catalyst Optimizer Phase – Analysis

- **Transform** an “**unresolved logical plan**” (tree build on an input DataFrame, DataSet or SQL AST) with unresolved attributes references and data types into a “**resolved logical plan**”
- Applies rules for:
 - Looking up relations by name from the catalog
 - Mapping named attributes to the input provided given operator’s children.
 - Determining which attributes refer to the same value to give them a unique ID
 - Propagating and coercing types through expressions:
 - for example, we cannot know the return type of $1 + col$ until we have resolved col and possibly casted its subexpressions to a compatible types.
- An attribute is called unresolved if we do not know its type or have not matched it to an input table (or an alias) → will fail into an error and exit
- Catalyst rules / transforms and Catalog object are used to track the tables in all data sources and resolve attributes

A Query Example – Logical Plan

- A Logical Plan describes a series of computation on datasets without defining the execution details
- It includes / provides:
 - outputs : a list of attributes generated by the logical plan (e.g. [sum] or [id, v])
 - constraints: a set of invariants about the result generated by the plan (e.g. $t2.id > 50000$) not to be confused with conditions like $t1.id = t2.id$
 - statistics: size of the plan in row/bytes + per columns stats (min/max/nulls/number of distinct values)

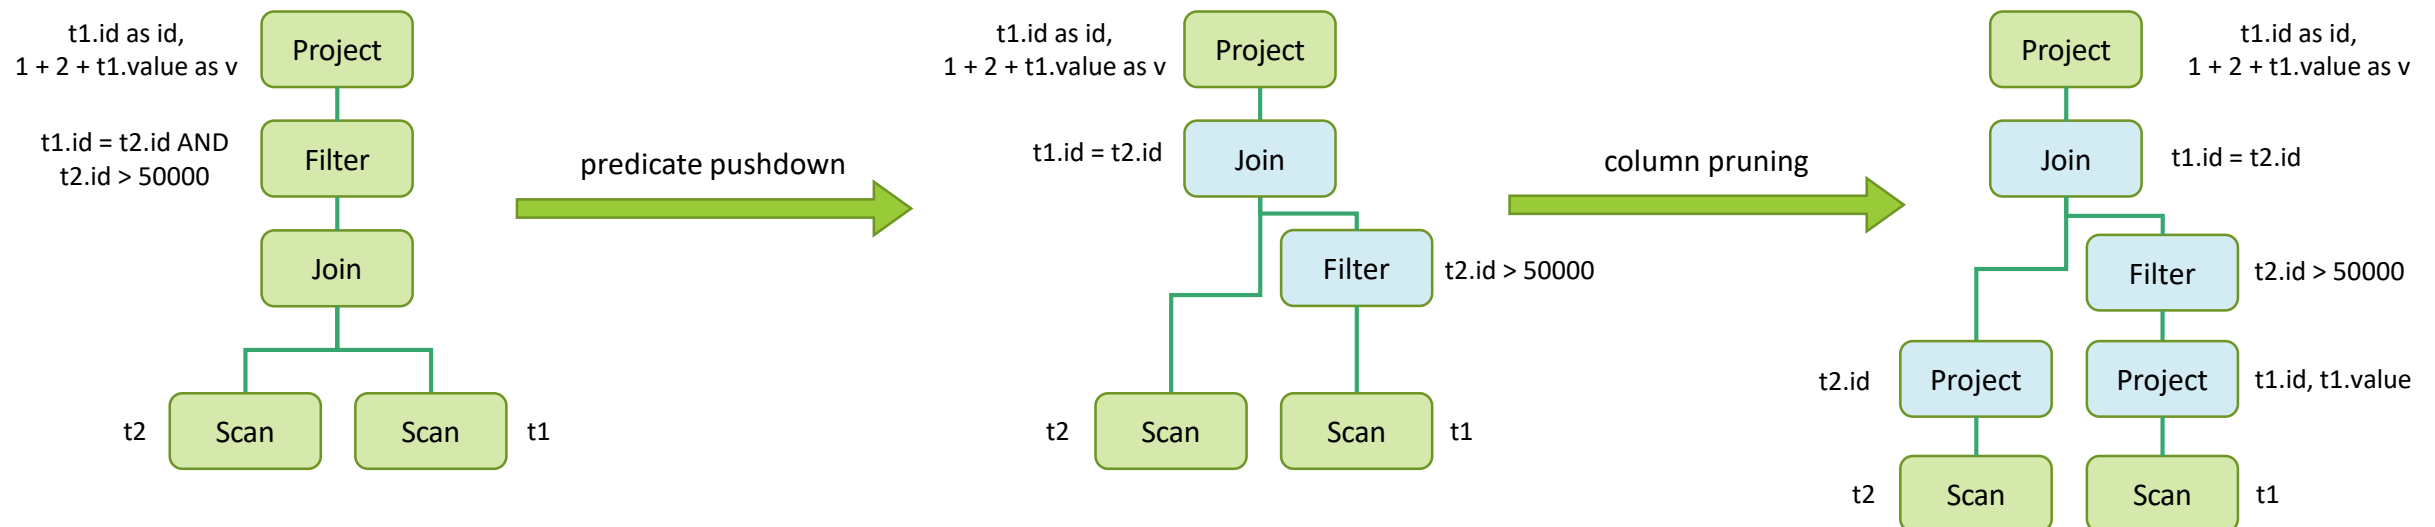


Catalyst Optimizer Phase

Logical Optimization

Catalyst Optimizer Phase (2/4) – Logical Optimization

- Applies standard rule-based optimizations to a logical plan
- Cost-based optimization is performed by generating multiple plans using rules, and then computing their costs
- These optimizations include:
 - constant folding
 - predicate pushdown
 - projection pruning
 - null propagation
 - boolean expression simplification
 - and many other rules...



Catalyst Optimizer Phase

Physical Planning

Catalyst Optimizer Phase (3/4) – Physical Planning

- Use a physical planner in phases:
 - Uses Strategies with pattern matching to generate one or more Physical Plans from an Optimized Logical Plan
 - Then, apply Rules to make the Physical Plans ready for execution:
 - Prepare Scalar sub-queries
 - Ensure requirements on input rows (partitioning for example)
 - Apply physical optimizations (rule-based, such as pipelining projections or filters)
- Then, selects a plan using a cost model


Catalyst Optimizer Phase

Code Generation

Catalyst Optimizer Phase (4/4) – Code Generation

- The goal is to generate Java bytecode to run on each Spark node
- Initially, relied on a Scala feature: “quasiquotes”
 - allow the construction of abstract syntax trees (ASTs) in the Scala language
 - can then be fed to the Scala compiler at runtime to generate bytecode
- Now uses other framework named Janino

- Code with performance comparison:



- Converts an expression like $(x+y) + 1$ to:

```
def compile (node: Node) : AST = node match {  
  case Literal (value) => q"$value"  
  case Attribute (name) => q"row.get($name)"  
  case Add(left, right) => q"{compile($left)} + {compile($right)}"  
}
```

- 100+ native functions with optimized code generation implementations:
 - String : concat, format_string, lower, lpad
 - Data/Time : current_timestamp, date_format, date_add
 - Math : sqrt, randn